



SUMMIT
ONLINE



Modernize your data warehouse

Aneesh Chandra PN

Data & Analytics Specialist Solutions Architect

Amazon Web Services



Agenda

- Modern analytics and Amazon Redshift
- Architecture and concepts
- Accelerating your data warehouse migration
- Additional resources
- Q&A

Data warehousing trends



Exponential growth
of event data



End-to-end insights from
analyzing all of your data



Migrations
to the cloud

Why Amazon Redshift


Tens of thousands of customers use Amazon Redshift and process over 2 EB of data per day



Data lake & AWS integrated

Lake formation catalog & security

Exabyte querying & AWS integrated
(e.g., AWS DMS, Amazon CloudWatch)




Best performance

3x faster than other
cloud data warehouses



Best value

Usage-based, RIs
Predictable costs



Most scalable

Virtually unlimited
elastic linear scaling



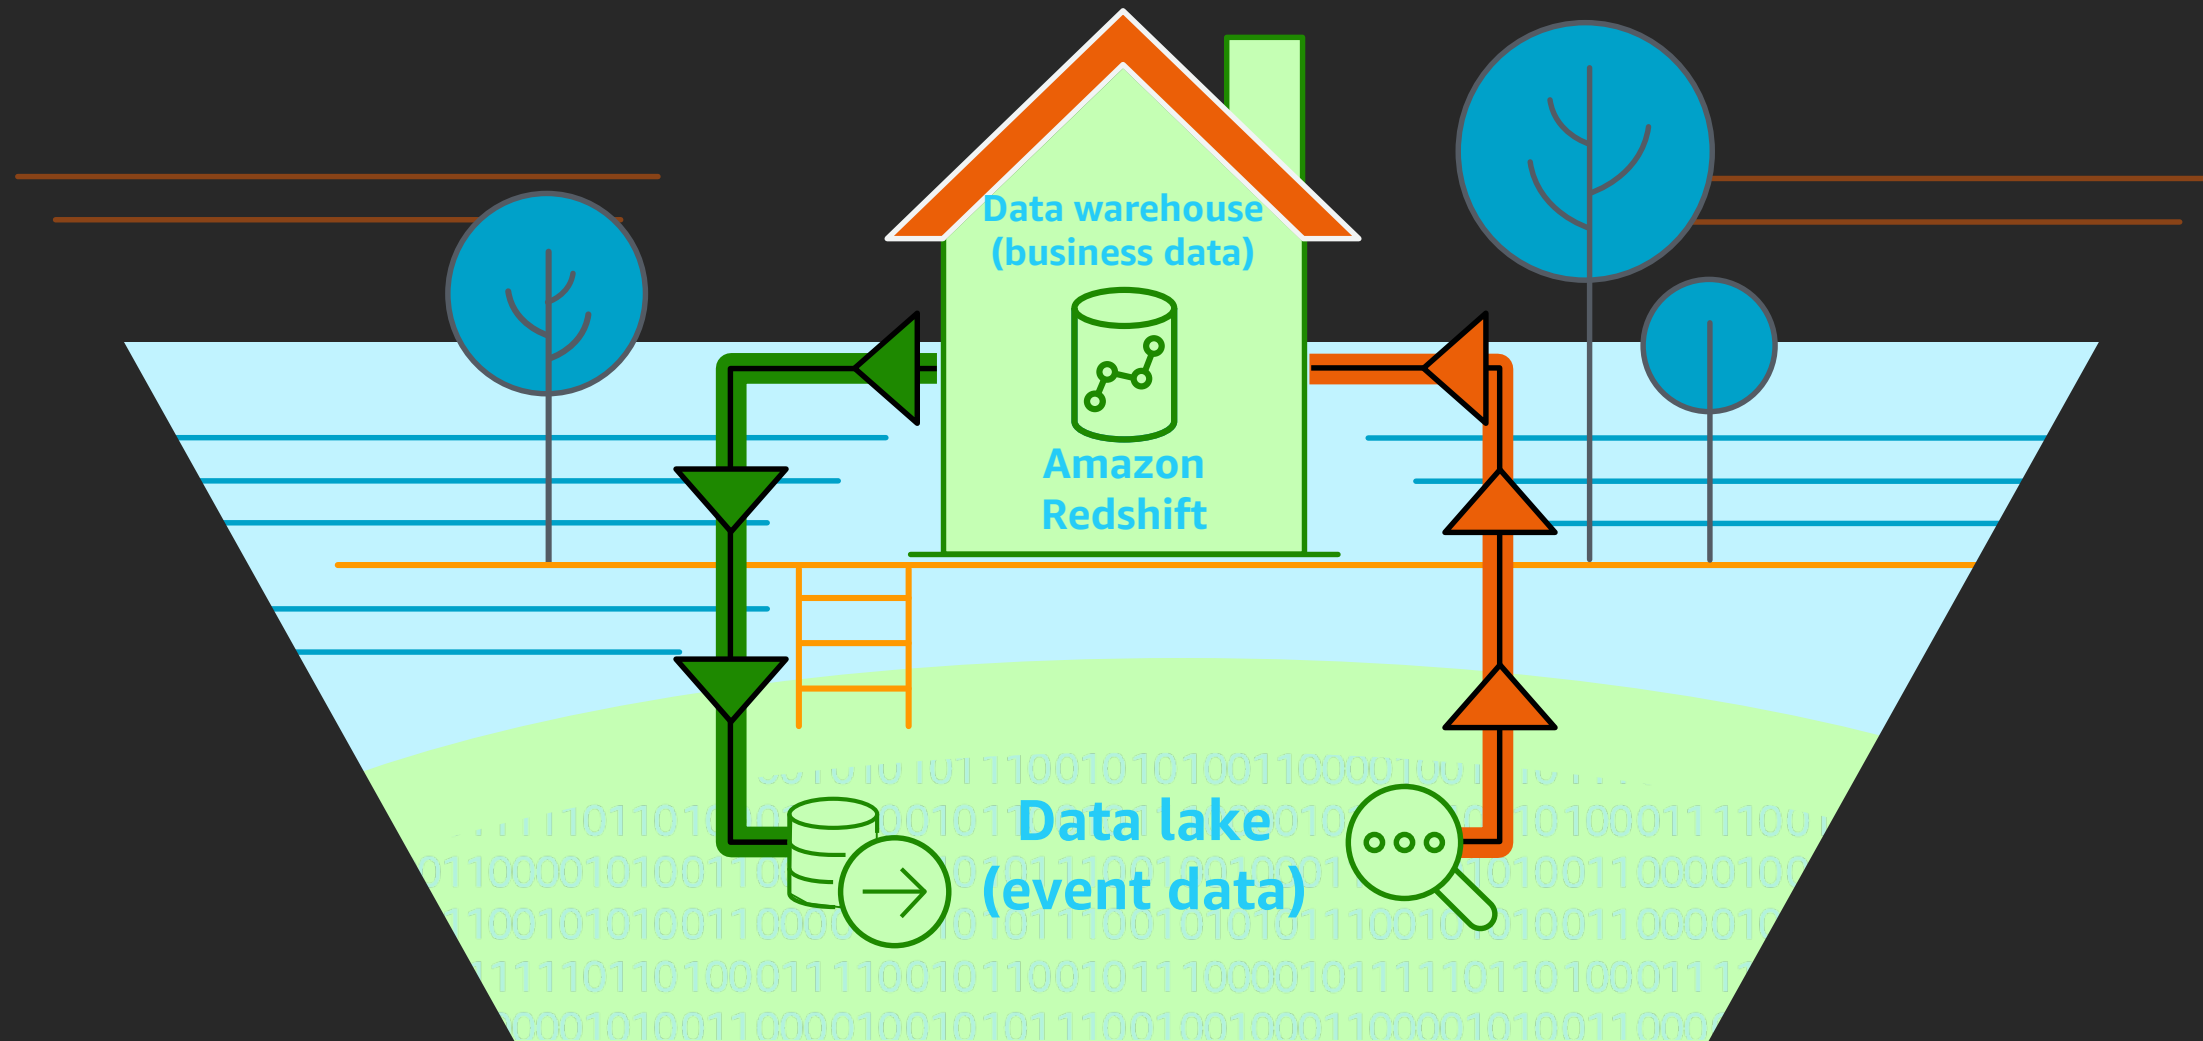
Most secure & compliant

AWS-grade security (e.g., VPC, encryption
with AWS KMS, AWS CloudTrail)



Easy to manage

Easy to provision & manage, automated
backups, AWS support, and 99.9% SLAs



Customers moving to data lake architectures

Amazon Redshift enables you to have a lake house approach




EQUINOX

“Moving to Amazon Redshift has helped us reduce our overall maintenance costs by nearly 80% compared with our legacy data warehouse. By leveraging Amazon Redshift Spectrum’s ability to query data directly in our Amazon S3 data lake, we have been able to easily integrate new data sources in hours, not days or weeks. This has not only reduced our time to insight, but it helped us control our infrastructure costs. Amazon Redshift requires very little maintenance, to the point where we don’t even have a dedicated administrator, and we spend less than an hour a month on maintenance and administration.”

Elliott Cordo
VP of Data Analytics
Equinox



Tens of thousands of customers use Amazon Redshift







Amazon Redshift – evolving architecture

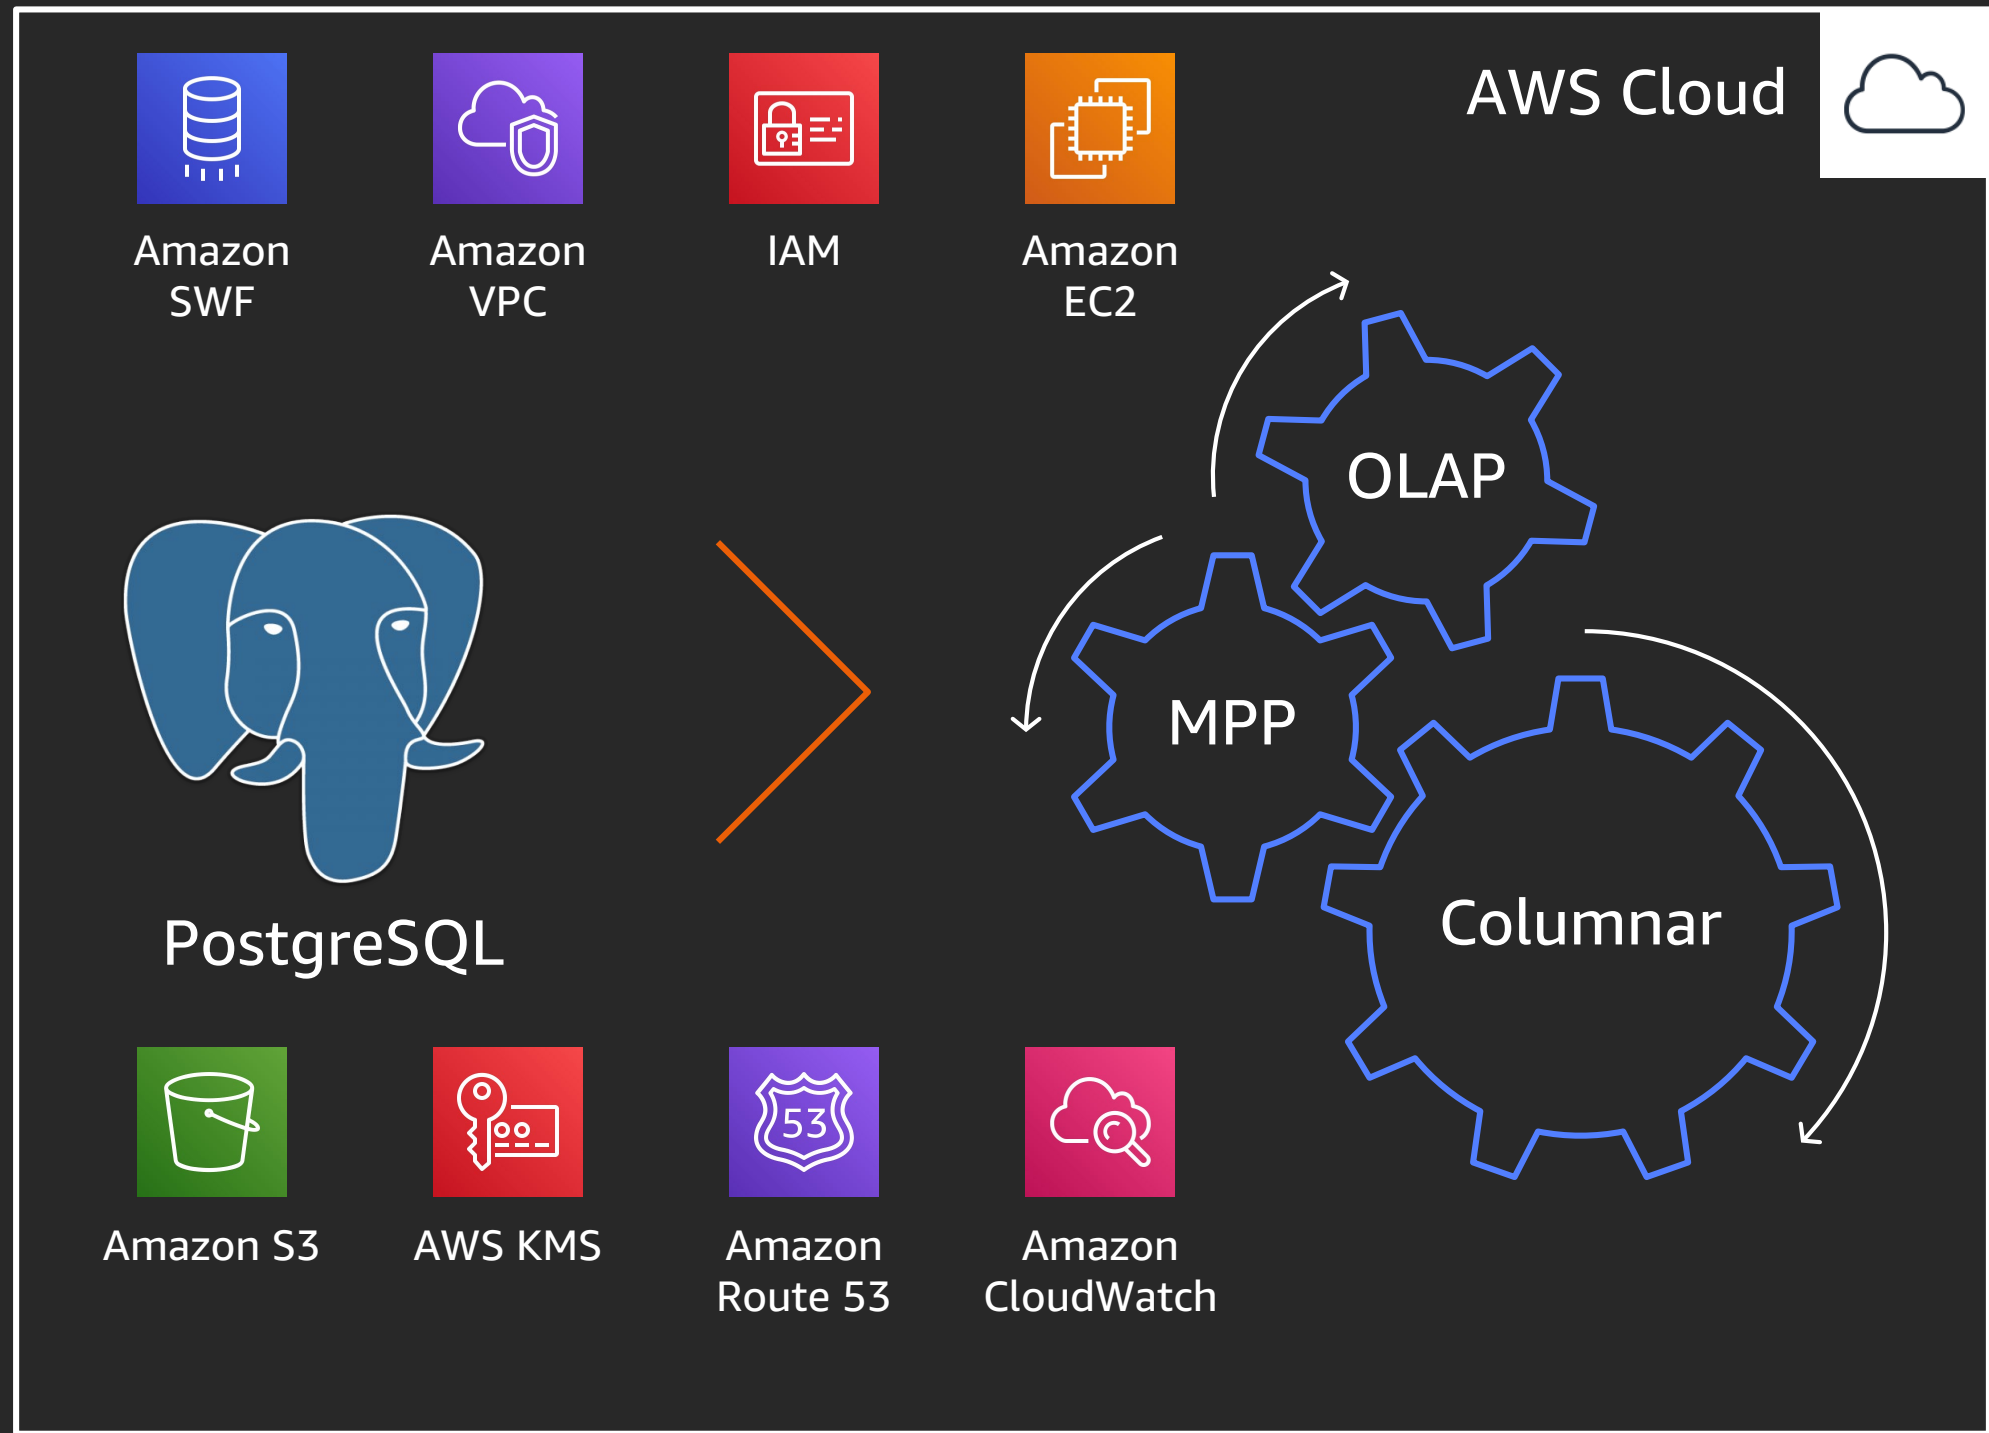


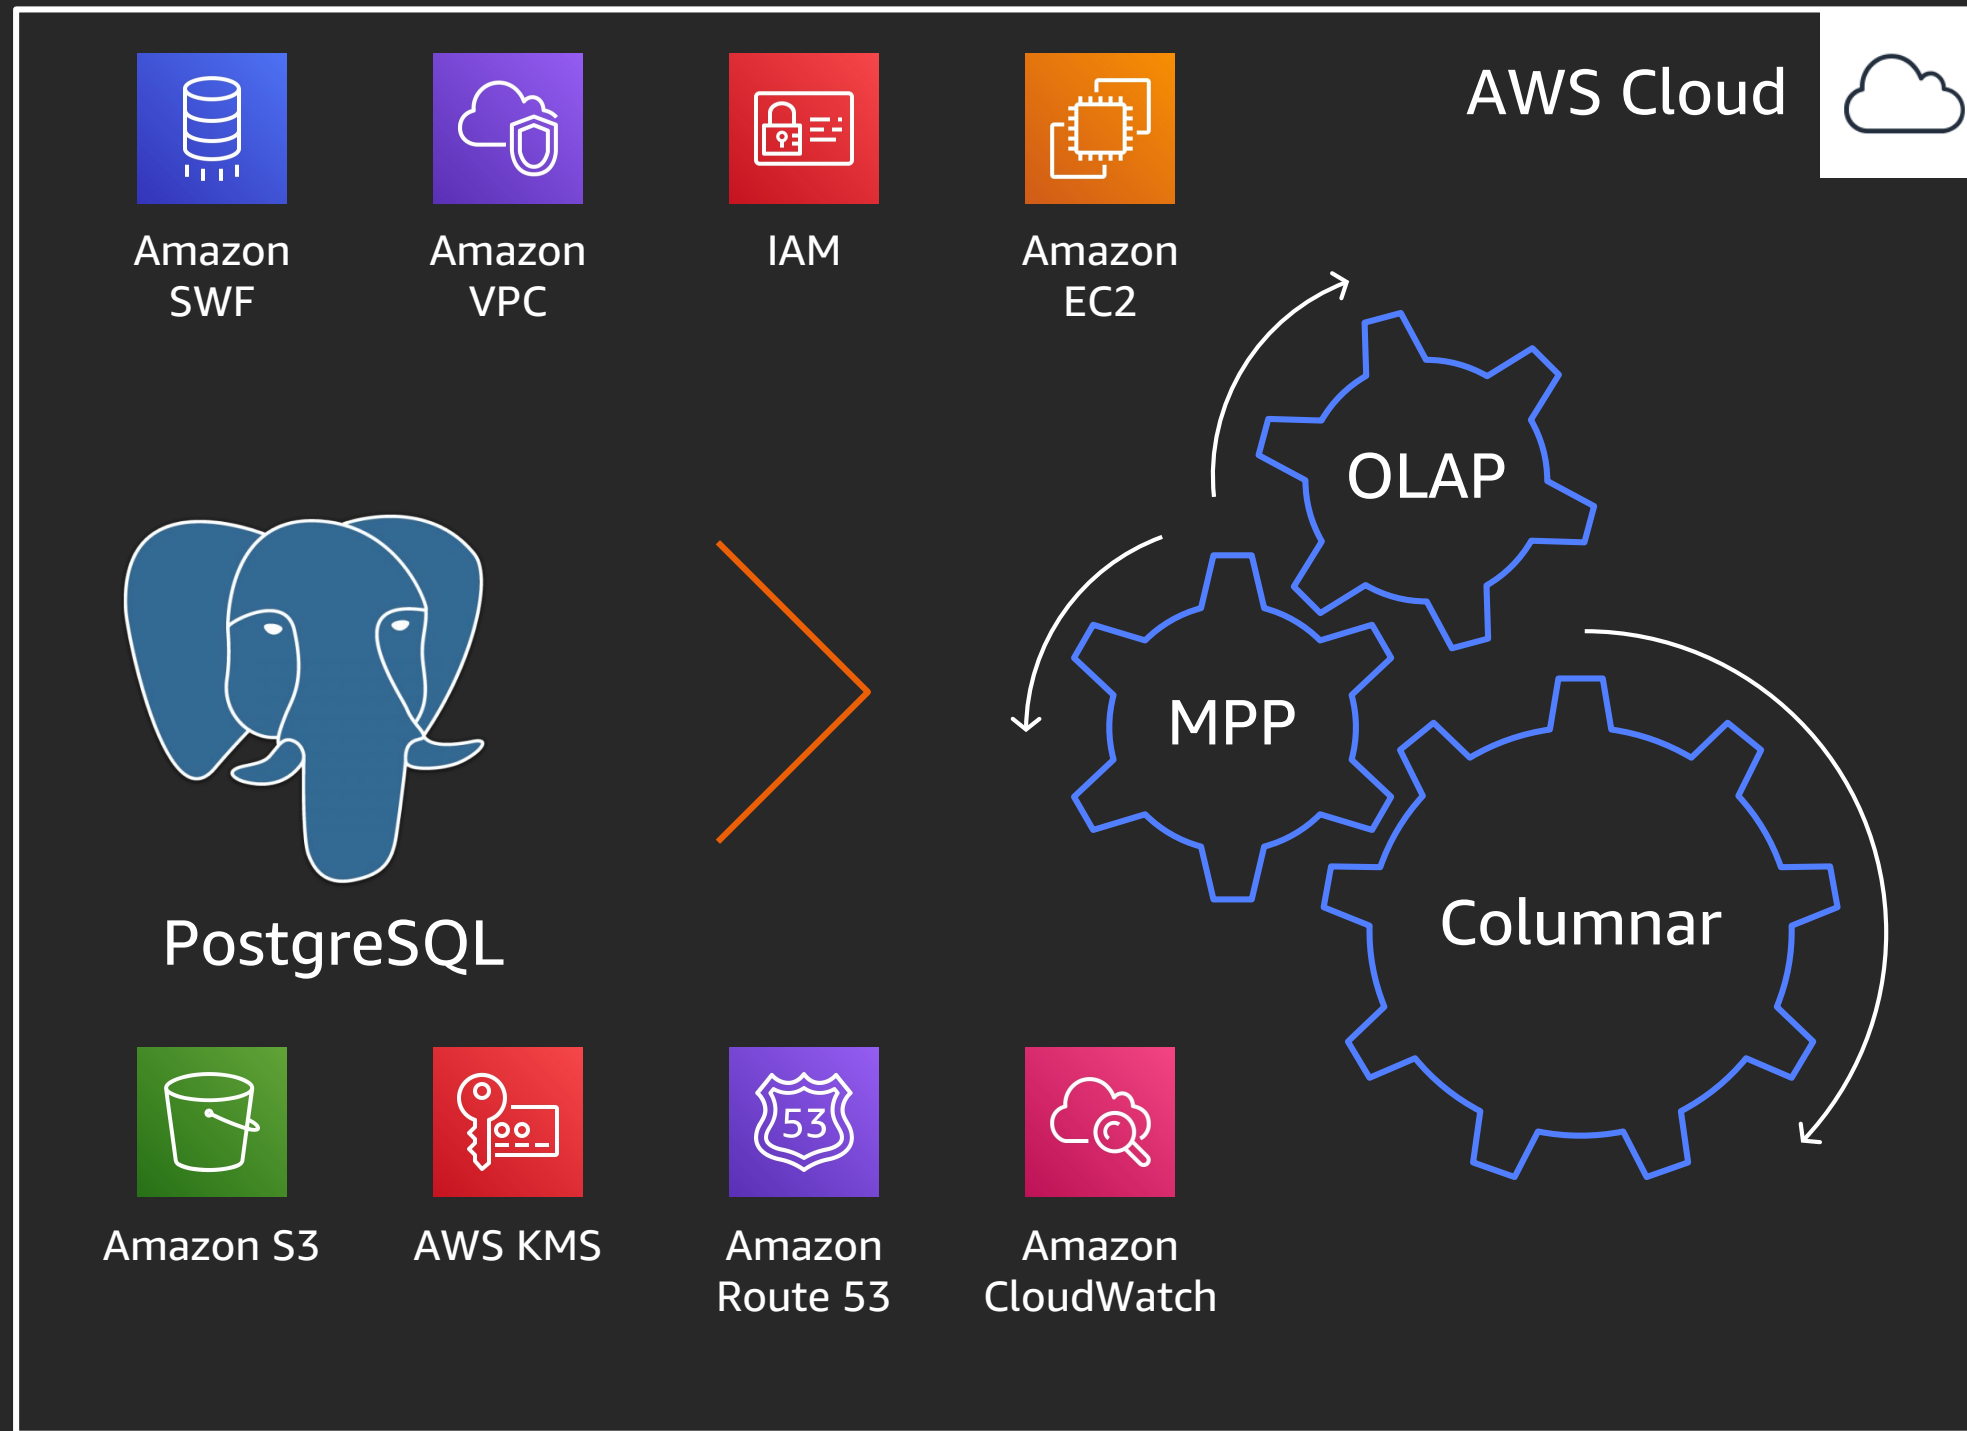
PostgreSQL



PostgreSQL







Features delivered to meet customer needs

Robust result set caching

Amazon Redshift Spectrum: date formats, scalar JSON and ION file format support, region expansion, predicate filtering

Unload to CSV

Auto WLM

Concurrency scaling

Manage multi-part query in AWS console

Amazon Redshift Spectrum: Row group filtering in Parquet and ORC, nested data support, enhanced VPC routing, multiple partitions

Auto WLM with query priorities

Spatial processing

Large # of tables support ~20,000

Auto-analyze

~25 query monitoring rules (QMR) support

Autoanalyze for incremental changes on table

Faster classic resize with optimized data transfer protocol

Snapshot scheduler

Stored procedures

Column-level access control with AWS Lake Formation

Copy command support for ORC, Parquet

Health and performance monitoring with Amazon CloudWatch

RA3

IAM role chaining

Automatic table distribution style

Advisor recommendations for distribution keys

Performance of inter-region snapshot transfers

Federated query

Elastic resize

CloudWatch support for WLM queues

AQUA (Advanced Query Accelerator)

DC1 migration to DC2

Spectrum Request Accelerator

Performance: Bloom filters in joins, complex queries that create internal table, communication layer

Integration with AWS Lake Formation

AZ64 compression encoding

Materialized views

Groups

Performance enhancements: hash join, vacuum, window functions, resize ops, aggregations, console, union all, efficient compile code cache

Resiliency of ROLLBACK processing

Apply new distribution key

Amazon Redshift Spectrum: Concurrency scaling

Autovacuum sort, autoanalyze, and autotable sort

Console redesign

Manual pause and resume

200+

new features in the past 18 months

Amazon Redshift architecture

Massively parallel, shared-nothing
columnar architecture

Leader node

SQL endpoint

Stores metadata

Coordinates parallel SQL processing

Compute nodes

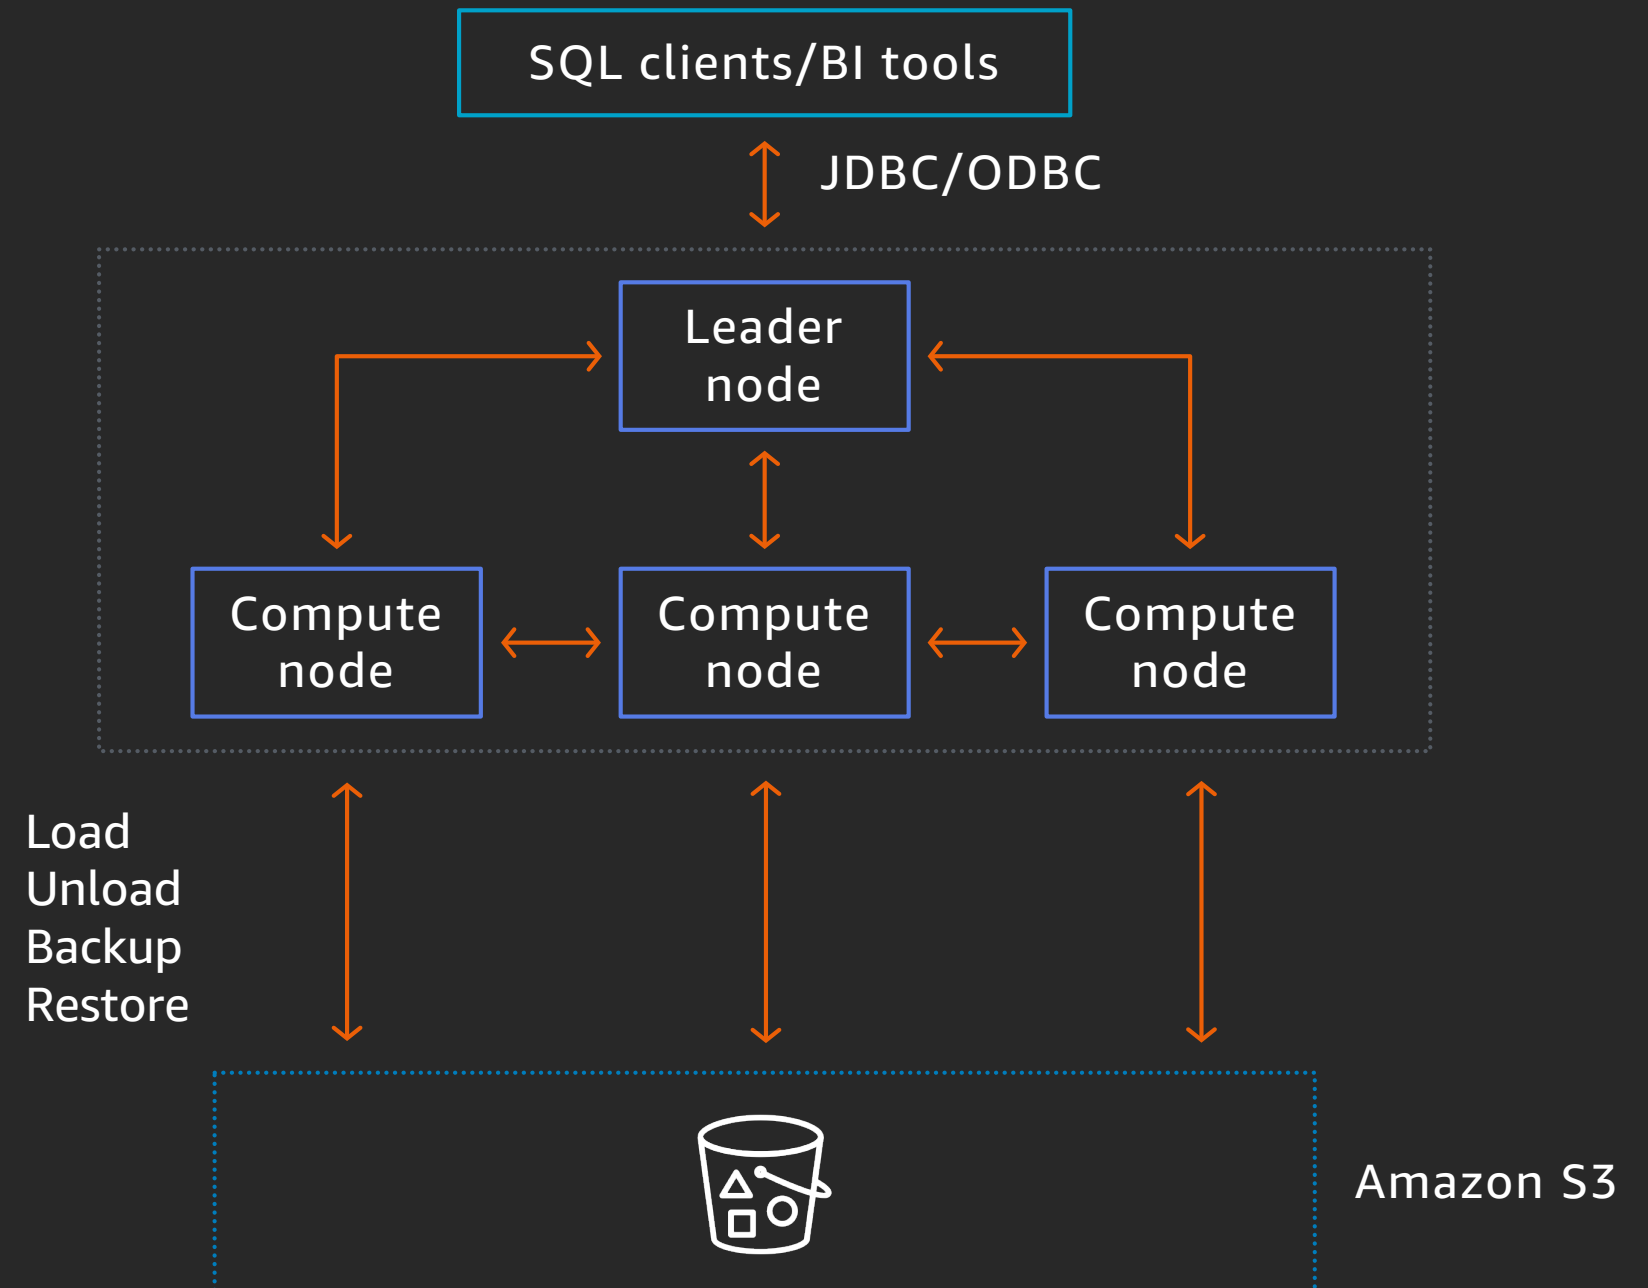
Local, columnar storage

Executes queries in parallel

Load, unload, backup, restore

Amazon Redshift Spectrum nodes

Execute queries directly against Amazon S3



Amazon Redshift architecture

Massively parallel, shared-nothing columnar architecture

Leader node

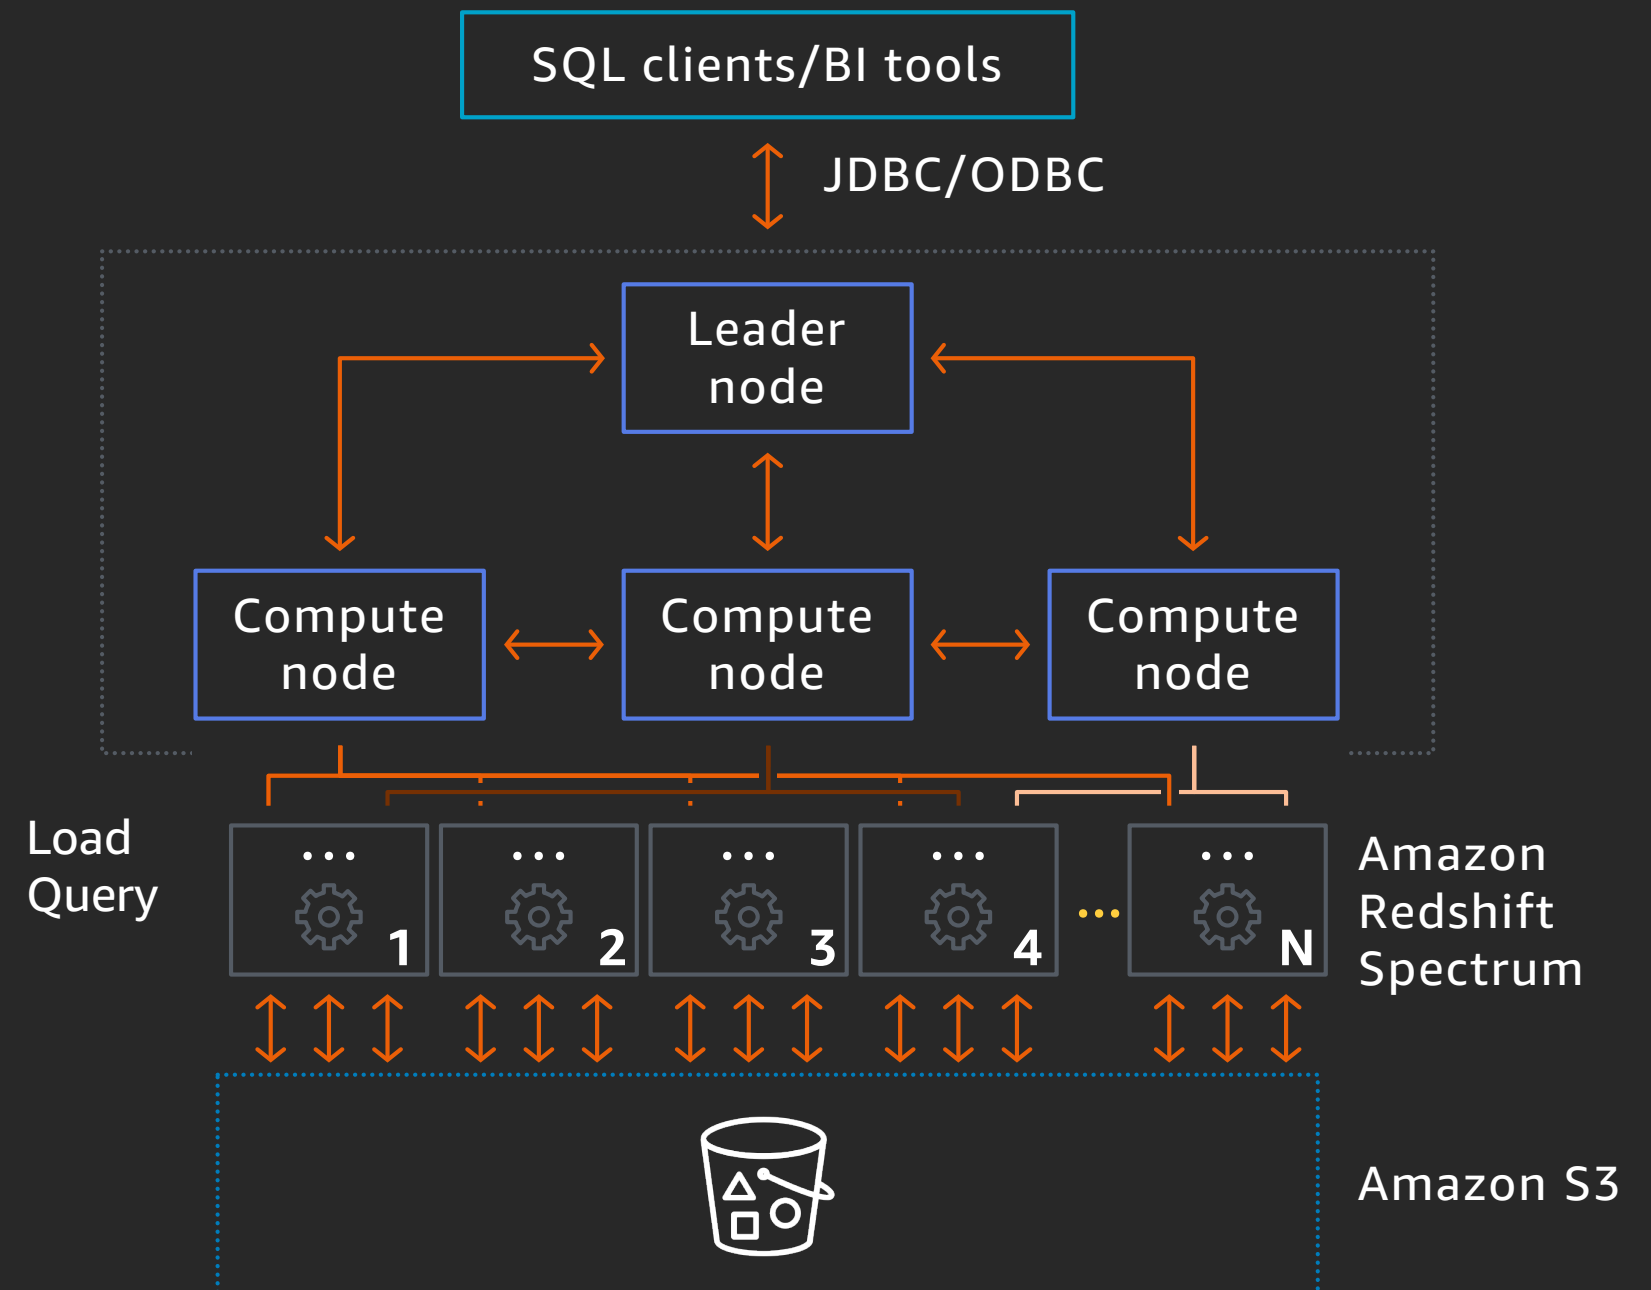
- SQL endpoint
- Stores metadata
- Coordinates parallel SQL processing

Compute nodes

- Local, columnar storage
- Executes queries in parallel
- Load, unload, backup, restore

Amazon Redshift Spectrum nodes

- Execute queries directly against Amazon S3



Amazon Redshift evolving architecture

Massively parallel, shared-nothing
columnar architecture

Leader node

Compute nodes

Amazon Redshift Spectrum nodes

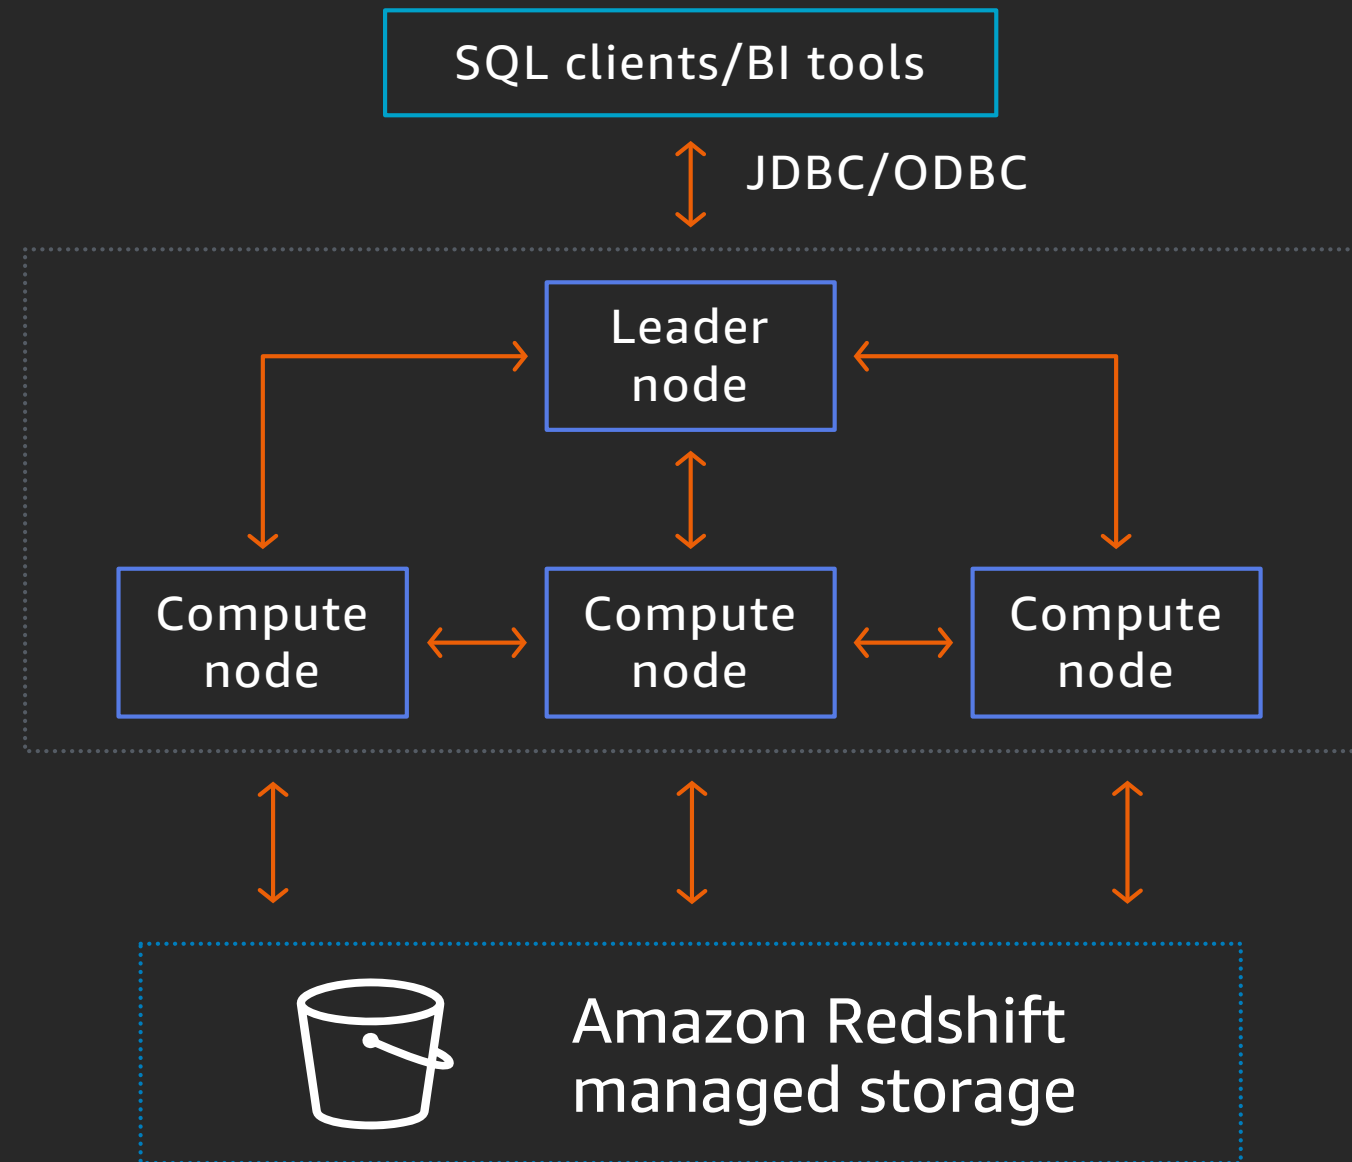
Amazon Redshift managed storage

Pay separately for storage and compute

Large high-speed SSD-backed cache

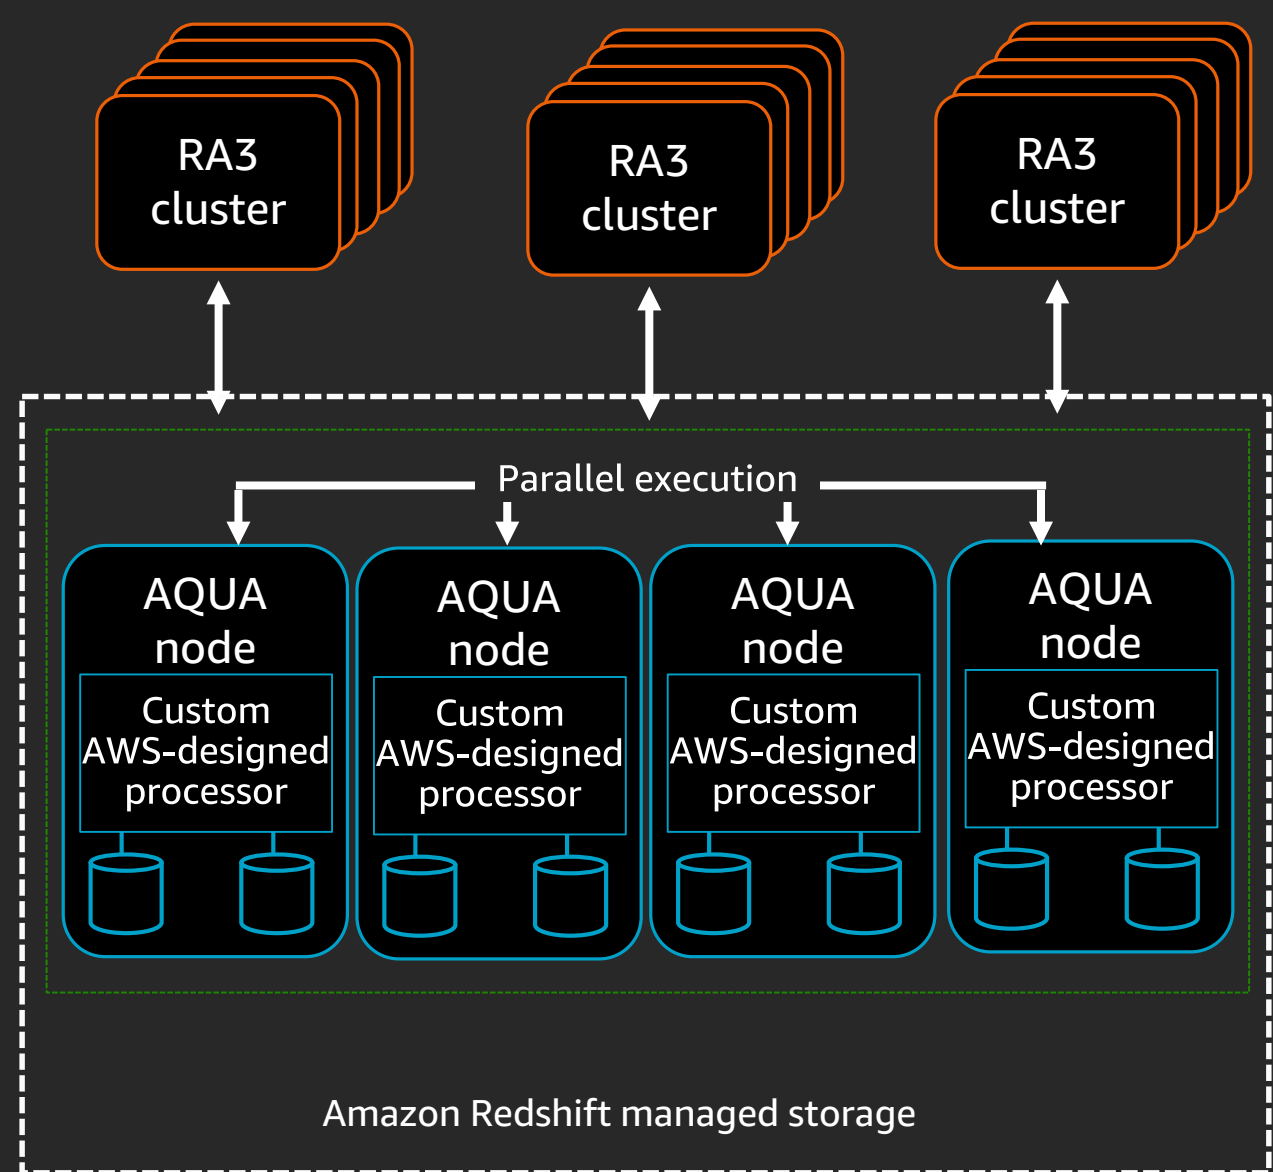
Automatic scaling (up to 64 TB/instance)

Supports up to 8.2 PB of cluster storage



AQUA: Advanced Query Accelerator (preview)

A new distributed and hardware-accelerated processing layer that will make Amazon Redshift 10x faster than any other cloud data warehouse without increasing cost



Minimize data movement over the network by pushing down operations to AQUA nodes

AQUA nodes with custom AWS-designed analytics processors to make operations (compression, encryption, filtering, and aggregations) faster than traditional CPUs

Available in preview only with RA3 – no code changes required



Terminology and concepts



Terminology and concepts: Node types

Amazon Redshift analytics – RA3 (new)

Amazon Redshift managed storage – solid-state disks + Amazon S3

Dense compute – DC2

Solid-state disks

Dense storage – DS2

Magnetic disks

Instance type	Disk type	Size	Memory	CPUs	Slices
RA3 4xlarge (new)	RMS	Scales to 16 TB	96 GB	12	4
RA3 16xlarge (new)	RMS	Scales to 64 TB	384 GB	48	16
DC2 large	SSD	160 GB	16 GB	2	2
DC2 8xlarge	SSD	2.56 TB	244 GB	32	16
DS2 xlarge	Magnetic	2 TB	32 GB	4	2
DS2 8xlarge	Magnetic	16 TB	244 GB	36	16



Terminology and concepts: Columnar

Amazon Redshift uses a columnar architecture for storing data on disk

Goal: Reduce I/O for analytics queries

Physically store data on disk by column rather than row

Read only the column data that is required

Example: Columnar architecture

```
CREATE TABLE deep_dive (  
    aid INT      --audience_id  
    ,loc CHAR(3) --location  
    ,dt  DATE    --date  
);
```

aid	loc	dt

aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

```
SELECT min(dt) FROM deep_dive;
```

Row-based storage

- Need to read everything
- Unnecessary I/O

Example: Columnar architecture

```
CREATE TABLE deep_dive (  
    aid INT      --audience_id  
    ,loc CHAR(3) --location  
    ,dt  DATE    --date  
);
```

aid	loc	dt

aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

```
SELECT min(dt) FROM deep_dive;
```

Column-based storage

Only scan blocks for relevant column



Terminology and concepts: Compression

Goals

Allow more data to be stored within an Amazon Redshift cluster

Improve query performance by decreasing I/O

Impact

Allows 2–4x more data to be stored within the cluster

By default, COPY automatically analyzes and compresses data on first load into an empty table

ANALYZE COMPRESSION is a built-in command that finds the optimal compression for each column on an existing table

Compression example

```
CREATE TABLE deep_dive (  
  aid INT      --audience_id  
  ,loc CHAR(3) --location  
  ,dt  DATE    --date  
);
```

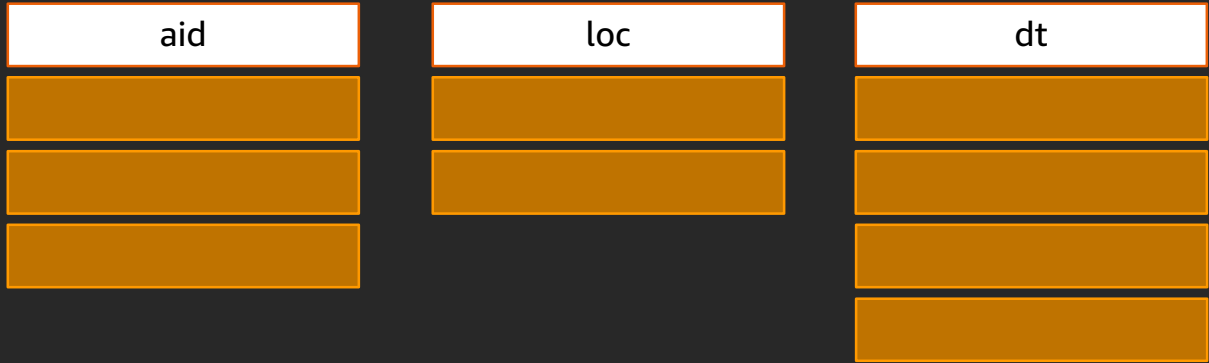
aid	loc	dt

aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

Add 1 of 13 different encodings to each column

Compression example

```
CREATE TABLE deep_dive (  
  aid INT          ENCODE AZ64  
  ,loc CHAR(3)     ENCODE BYTEDICT  
  ,dt  DATE        ENCODE RUNLENGTH  
);
```



aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

- More efficient compression by storing the same data type in the columnar architecture
- Columns grow and shrink independently
- Reduces storage requirements
- Reduces I/O



Terminology and concepts: Blocks

Column data is persisted to 1 MB immutable blocks

Blocks are individually encoded with 1 of 13 encodings

A full block can contain millions of values



Terminology and concepts: Zone maps

Goal

Eliminates unnecessary I/O

In-memory block metadata

- Contains per-block minimum and maximum values
- All blocks automatically have zone maps
- Effectively prunes blocks that cannot contain data for a given query



Terminology and concepts: Data sorting

Goal

Make queries run faster by increasing the effectiveness of zone maps and reducing I/O

Impact

Enables range-restricted scans to prune blocks by leveraging zone maps

Achieved with the table property SORTKEY defined on one or more columns

Optimal sort key is dependent on:

- Query patterns
- Business requirements
- Data profile

Example: Sort key

```
CREATE TABLE deep_dive (  
  aid INT --audience_id  
  ,loc CHAR(3) --location  
  ,dt DATE --date  
) SORTKEY (dt, loc);
```

Add a sort key to one or more columns to physically sort the data on disk

deep_dive		
aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

Example: Sort key

```
CREATE TABLE deep_dive (  
  aid INT --audience_id  
  ,loc CHAR(3) --location  
  ,dt DATE --date  
) SORTKEY (dt, loc);
```

Add a sort key to one or more columns to physically sort the data on disk

deep_dive		
aid	loc	dt
1	SFO	2017-10-20
2	JFK	2017-10-20
3	SFO	2017-04-01
4	JFK	2017-05-14

deep_dive (sorted)		
aid	loc	dt
3	SFO	2017-04-01
4	JFK	2017-05-14
2	JFK	2017-10-20
1	SFO	2017-10-20



Example: Zone maps and sorting

```
SELECT count(*) FROM deep_dive WHERE dt = '06-09-2017';
```

Unsorted table

	MIN: 01-JUNE-2017
	MAX: 20-JUNE-2017

	MIN: 08-JUNE-2017
	MAX: 30-JUNE-2017

	MIN: 12-JUNE-2017
	MAX: 20-JUNE-2017

	MIN: 02-JUNE-2017
	MAX: 25-JUNE-2017



Example: Zone maps and sorting

```
SELECT count(*) FROM deep_dive WHERE dt = '06-09-2017';
```

Unsorted table



MIN: 01-JUNE-2017
MAX: 20-JUNE-2017



MIN: 08-JUNE-2017
MAX: 30-JUNE-2017



MIN: 12-JUNE-2017
MAX: 20-JUNE-2017






MIN: 02-JUNE-2017
MAX: 25-JUNE-2017

Example: Zone maps and sorting

```
SELECT count(*) FROM deep_dive WHERE dt = '06-09-2017';
```

Unsorted table

	MIN: 01-JUNE-2017 MAX: 20-JUNE-2017
	MIN: 08-JUNE-2017 MAX: 30-JUNE-2017
	MIN: 12-JUNE-2017 MAX: 20-JUNE-2017
	MIN: 02-JUNE-2017 MAX: 25-JUNE-2017




Sorted by date

	MIN: 01-JUNE-2017 MAX: 06-JUNE-2017
	MIN: 07-JUNE-2017 MAX: 12-JUNE-2017
	MIN: 13-JUNE-2017 MAX: 21-JUNE-2017
	MIN: 21-JUNE-2017 MAX: 30-JUNE-2017


Example: Zone maps and sorting

```
SELECT count(*) FROM deep_dive WHERE dt = '06-09-2017';
```

Unsorted table

	MIN: 01-JUNE-2017 MAX: 20-JUNE-2017
	MIN: 08-JUNE-2017 MAX: 30-JUNE-2017
	MIN: 12-JUNE-2017 MAX: 20-JUNE-2017
	MIN: 02-JUNE-2017 MAX: 25-JUNE-2017

Sorted by date

	MIN: 01-JUNE-2017 MAX: 06-JUNE-2017
	MIN: 07-JUNE-2017 MAX: 12-JUNE-2017
	MIN: 13-JUNE-2017 MAX: 21-JUNE-2017
	MIN: 21-JUNE-2017 MAX: 30-JUNE-2017



Terminology and concepts: Slices

A slice can be thought of
as a virtual compute node

Unit of data partitioning
Parallel query processing

Facts about slices

Each compute node is initialized with
either 2 or 16 slices

Table rows are distributed to slices

A slice processes only its own data

Data distribution

Distribution style is a table property that dictates how that table's data is distributed throughout the cluster

KEY: Value is hashed, same value goes to same location (slice)

ALL: Full table data goes to the first slice of every node




EVEN: Round robin

AUTO: Combines EVEN and ALL

Goals

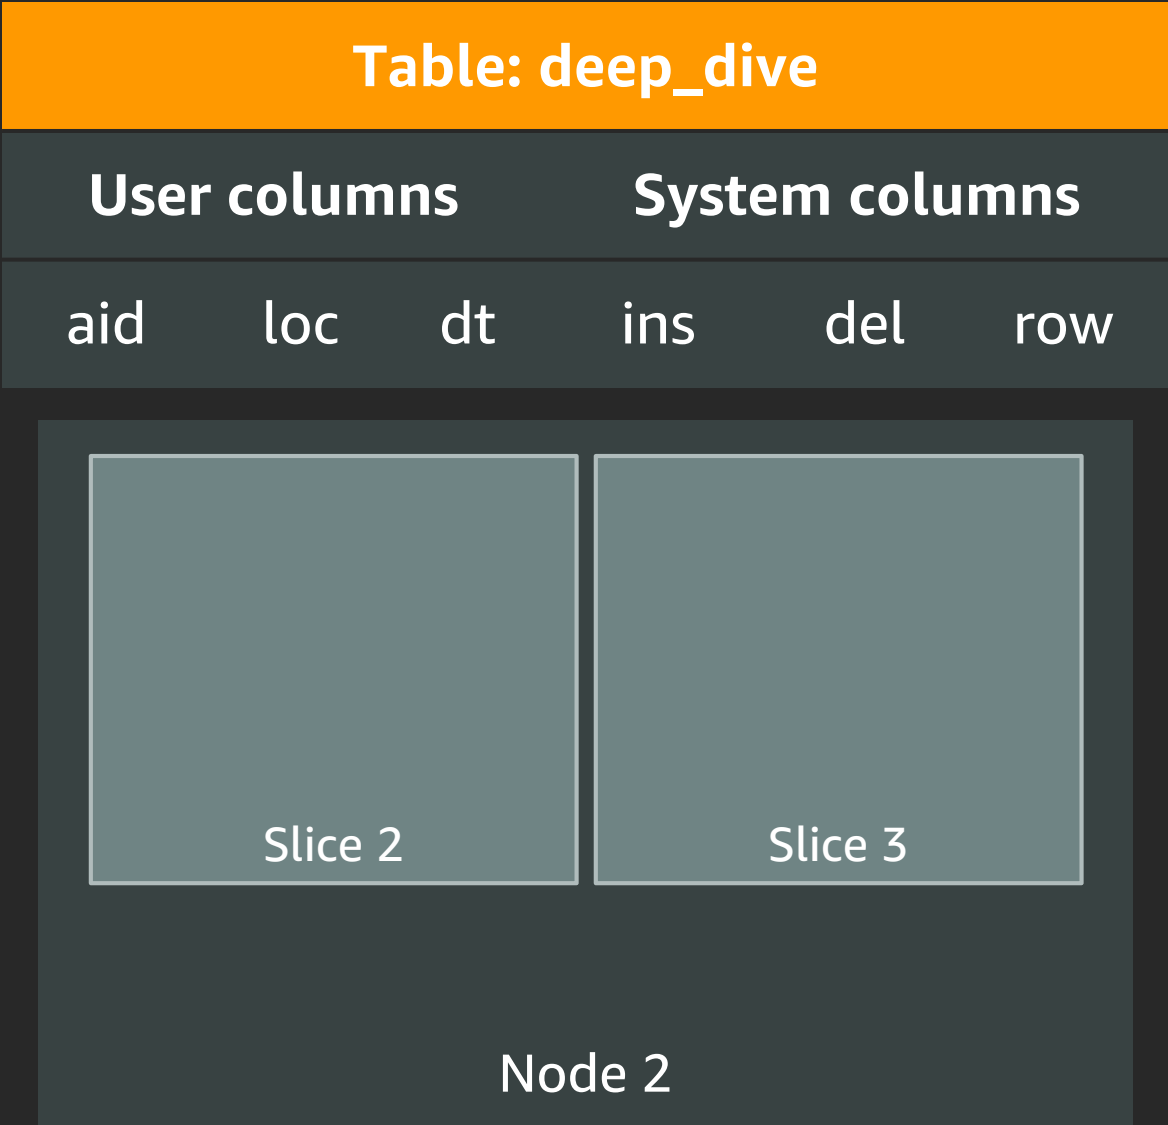
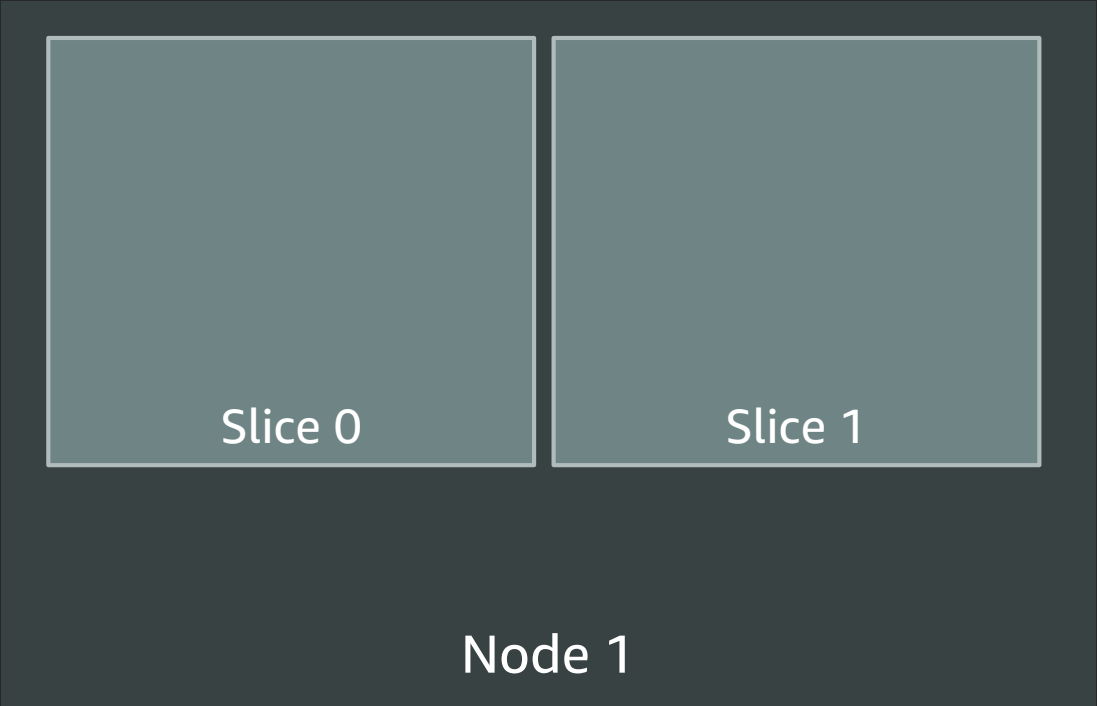
Distribute data evenly for parallel processing

Minimize data movement during query processing



Data distribution example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) (EVEN|KEY|ALL|AUTO);
```



Data distribution, **EVEN** example

```
CREATE TABLE deep_dive (  
    aid    INT        --audience_id  
    ,loc   CHAR(3)    --location  
    ,dt    DATE        --date  
) DISTSTYLE EVEN;
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row

Rows: 1

Slice 0

Slice 1

Node 1

Slice 2

Slice 3

Node 2

Data distribution, **EVEN** example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE EVEN;
```

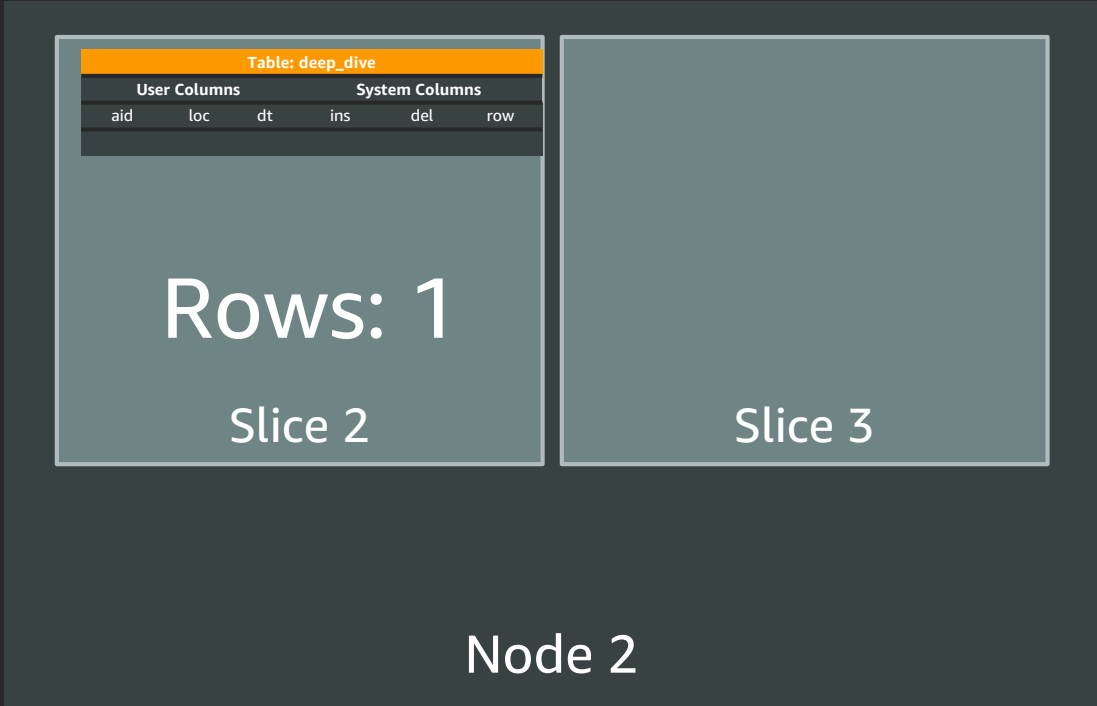
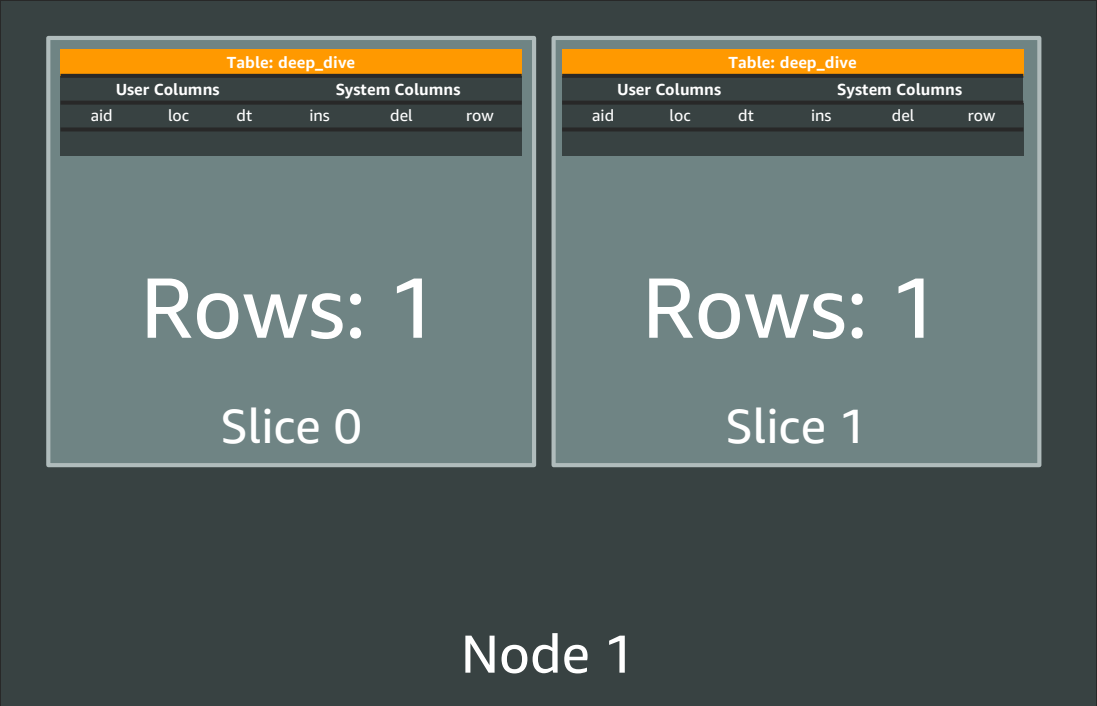
```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```



Data distribution, **EVEN** example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE EVEN;
```

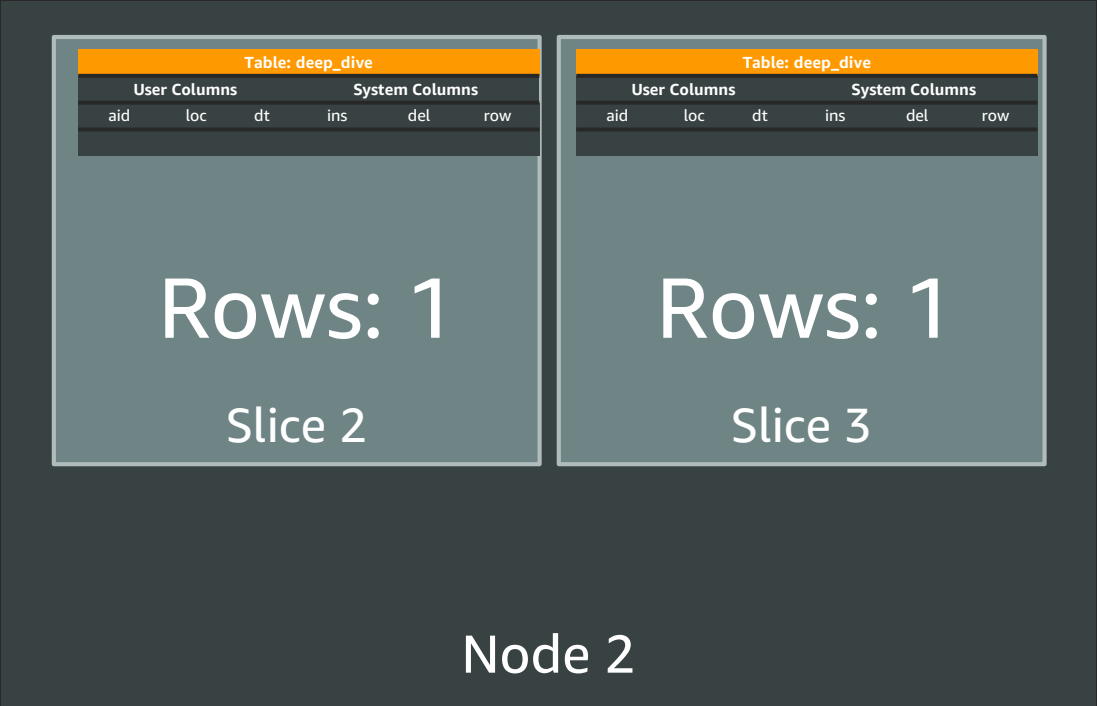

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```



Data distribution, **EVEN** example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE EVEN;
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```



Data distribution, **KEY** Example #1

```
CREATE TABLE deep_dive (  
    aid    INT        --audience_id  
    ,loc   CHAR(3)    --location  
    ,dt    DATE        --date  
) DISTSTYLE KEY DISTKEY (loc);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Rows: 0

Slice 0

Rows: 0

Slice 1

Node 1

Rows: 0

Slice 2

Rows: 0

Slice 3

Node 2

Data distribution, **KEY** Example #1

```
CREATE TABLE deep_dive (  
    aid    INT        --audience_id  
    ,loc   CHAR(3)    --location  
    ,dt    DATE        --date  
) DISTSTYLE KEY DISTKEY (loc);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row

Rows: 1

Slice 0

Rows: 0

Slice 1

Node 1

Rows: 0

Slice 2

Rows: 0

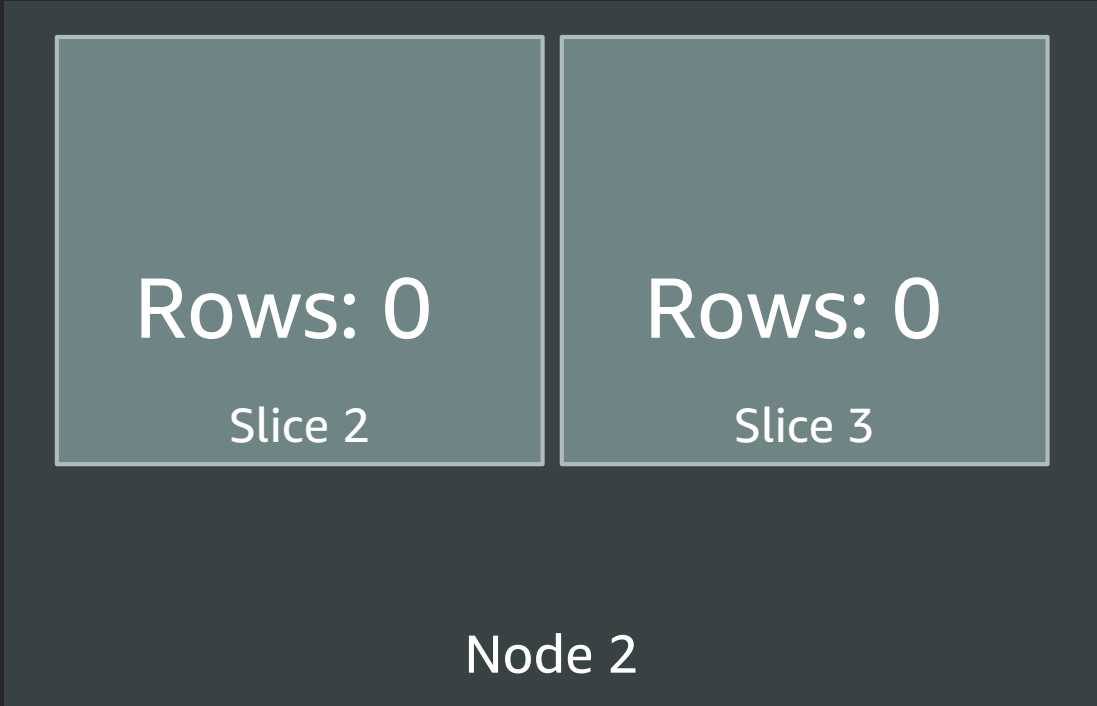
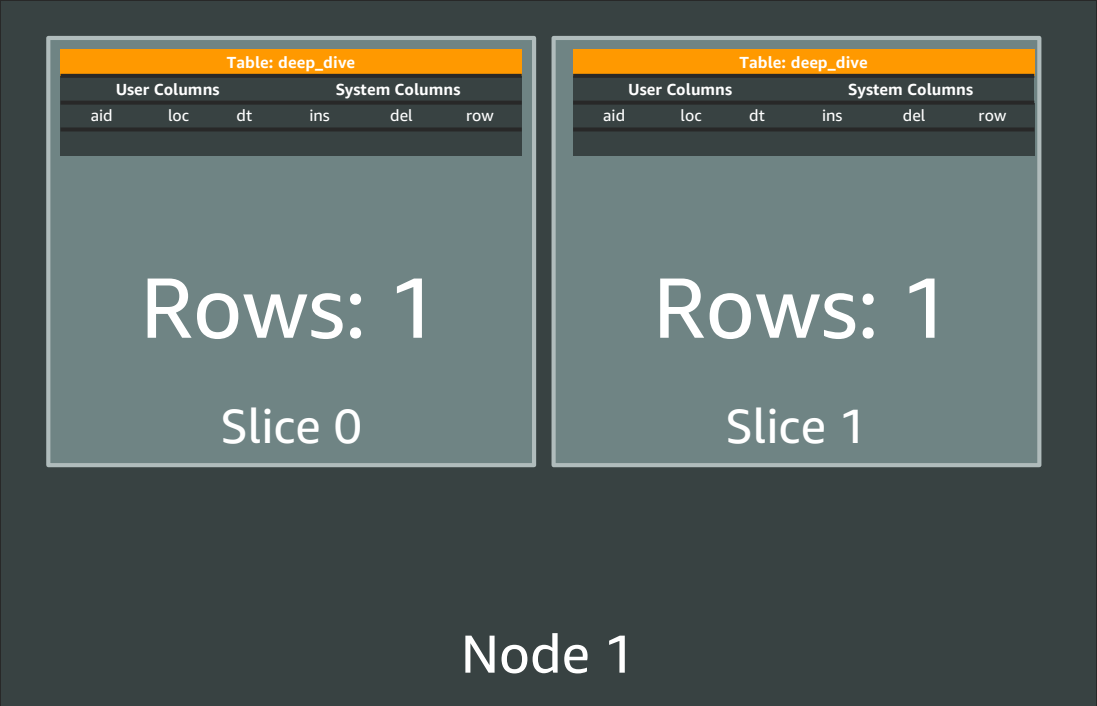
Slice 3

Node 2

Data distribution, **KEY** Example #1

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE KEY DISTKEY (loc);
```

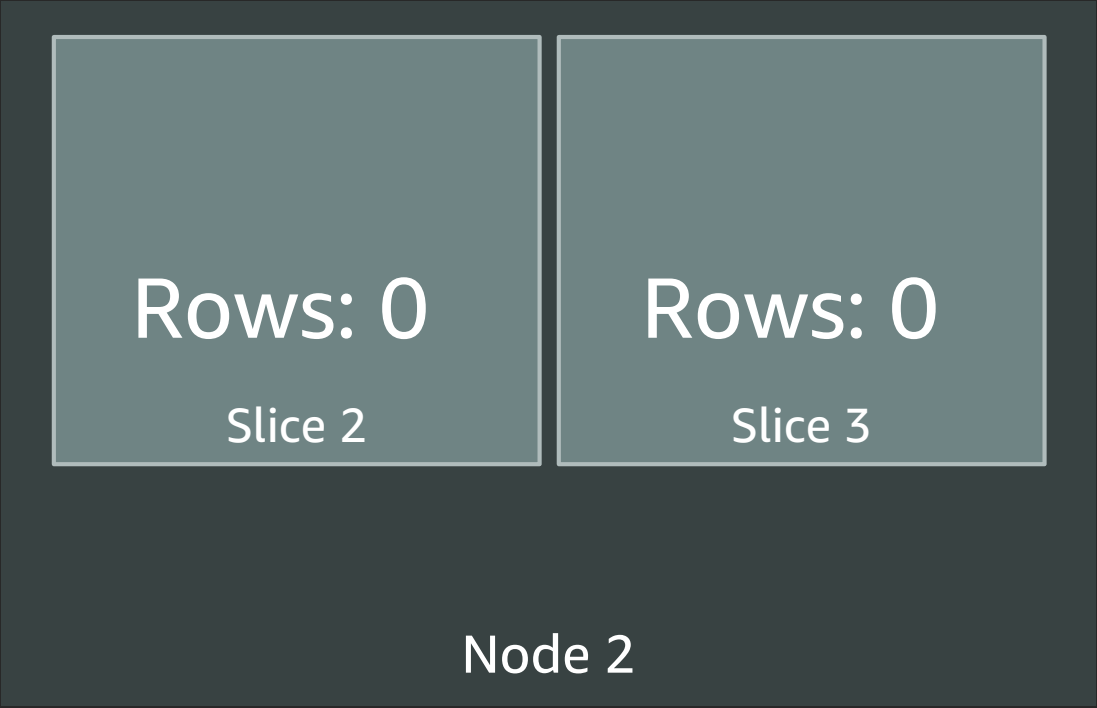
```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```



Data distribution, **KEY** Example #1

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE KEY DISTKEY (loc);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```



Data distribution, **KEY** Example #1

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE KEY DISTKEY (loc);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row

Rows: 2

Slice 0

Rows: 2

Slice 1

Node 1

Rows: 0

Slice 2

Rows: 0

Slice 3

Node 2

Data distribution, **KEY** Example #2

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE KEY DISTKEY (aid);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Rows: 0

Slice 0

Rows: 0

Slice 1

Node 1

Rows: 0

Slice 2

Rows: 0

Slice 3

Node 2

Data distribution, **KEY** Example #2

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE KEY DISTKEY (aid);
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 1					
Slice 0					

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 1					
Slice 1					

Node 1

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 1					
Slice 2					

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row
Rows: 1					
Slice 3					

Node 2

Data distribution, **ALL** example

```
CREATE TABLE deep_dive (  
  aid    INT      --audience_id  
  ,loc   CHAR(3)  --location  
  ,dt    DATE     --date  
) DISTSTYLE ALL;
```

```
INSERT INTO deep_dive VALUES  
(1, 'SFO', '2016-09-01'),  
(2, 'JFK', '2016-09-14'),  
(3, 'SFO', '2017-04-01'),  
(4, 'JFK', '2017-05-14');
```

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row

Rows: 4

Slice 0

Rows: 0

Slice 1

Node 1

Table: deep_dive					
User Columns			System Columns		
aid	loc	dt	ins	del	row

Rows: 4

Slice 2

Rows: 0

Slice 3

Node 2



Data distribution summary

DISTSTYLE KEY

Goals

- Optimize JOIN performance between large tables by distributing on columns used in the ON clause
- Optimize INSERT INTO SELECT performance
- Optimize GROUP BY performance

The column that is being distributed on should have a high cardinality and not cause row skew

```
SELECT diststyle, skew_rows
FROM svv_table_info WHERE "table" = 'deep_dive';
diststyle | skew_rows
-----+-----
KEY(aid) | 1.07
```

Ratio of the slice with the most and least number of rows

DISTSTYLE ALL

Goals

- Optimize JOIN performance with dimension tables
- Reduces disk usage on small tables

Small- and medium-size dimension tables (<3M rows)

DISTSTYLE EVEN

If neither KEY or ALL apply


DISTSTYLE AUTO

Default distribution combines DISTSTYLE ALL and EVEN



Accelerating your data warehouse migration

AWS migration tooling



AWS Schema Conversion Tool (AWS SCT) converts your commercial database and data warehouse schemas to open-source engines or AWS native services, such as Amazon Aurora and Amazon Redshift



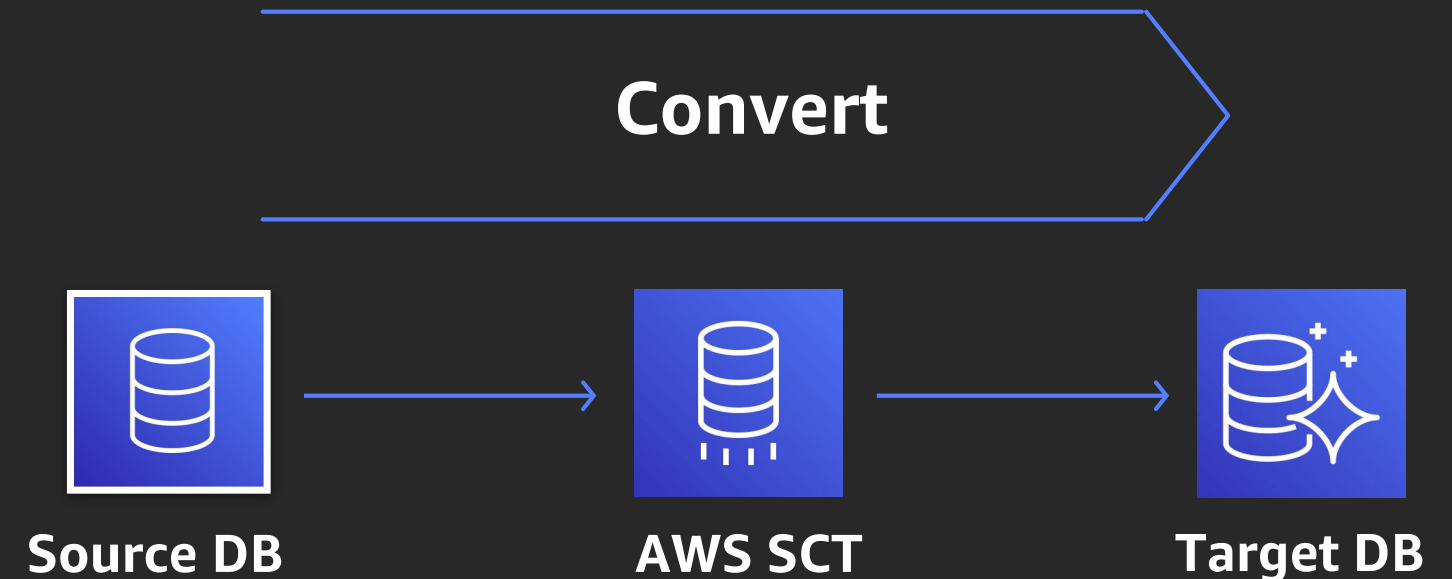
AWS Database Migration Service (AWS DMS) easily and securely migrates and/or replicates your databases and data warehouses to AWS

AWS SCT

AWS SCT helps automate database schema and code conversion tasks when migrating from source to target database engines

Features

- Create assessment reports for homogeneous/heterogeneous migrations
- Convert database schema
- Convert data warehouse schema
- Convert embedded application code
- Code browser that highlights places where manual edits are required
- Secure connections to your databases with SSL
- Service substitutions/ETL modernization to AWS Glue
- Migrate data to data warehouses using SCT data extractors
- Optimize schemas in Amazon Redshift



AWS SCT data extractors

Extract data from your data warehouse and migrate to Amazon Redshift

- **Extracts** data through local migration agents
- Data is **optimized** for Amazon Redshift and saved in local files
- Files are **loaded** to an Amazon S3 bucket (through network or AWS Snowball Edge) and then to Amazon Redshift


VERTICA

ORACLE

Microsoft SQL
Server


Greenplum

teradata.




AWS DMS



Migrating databases to AWS




Migrate between on-premises and AWS



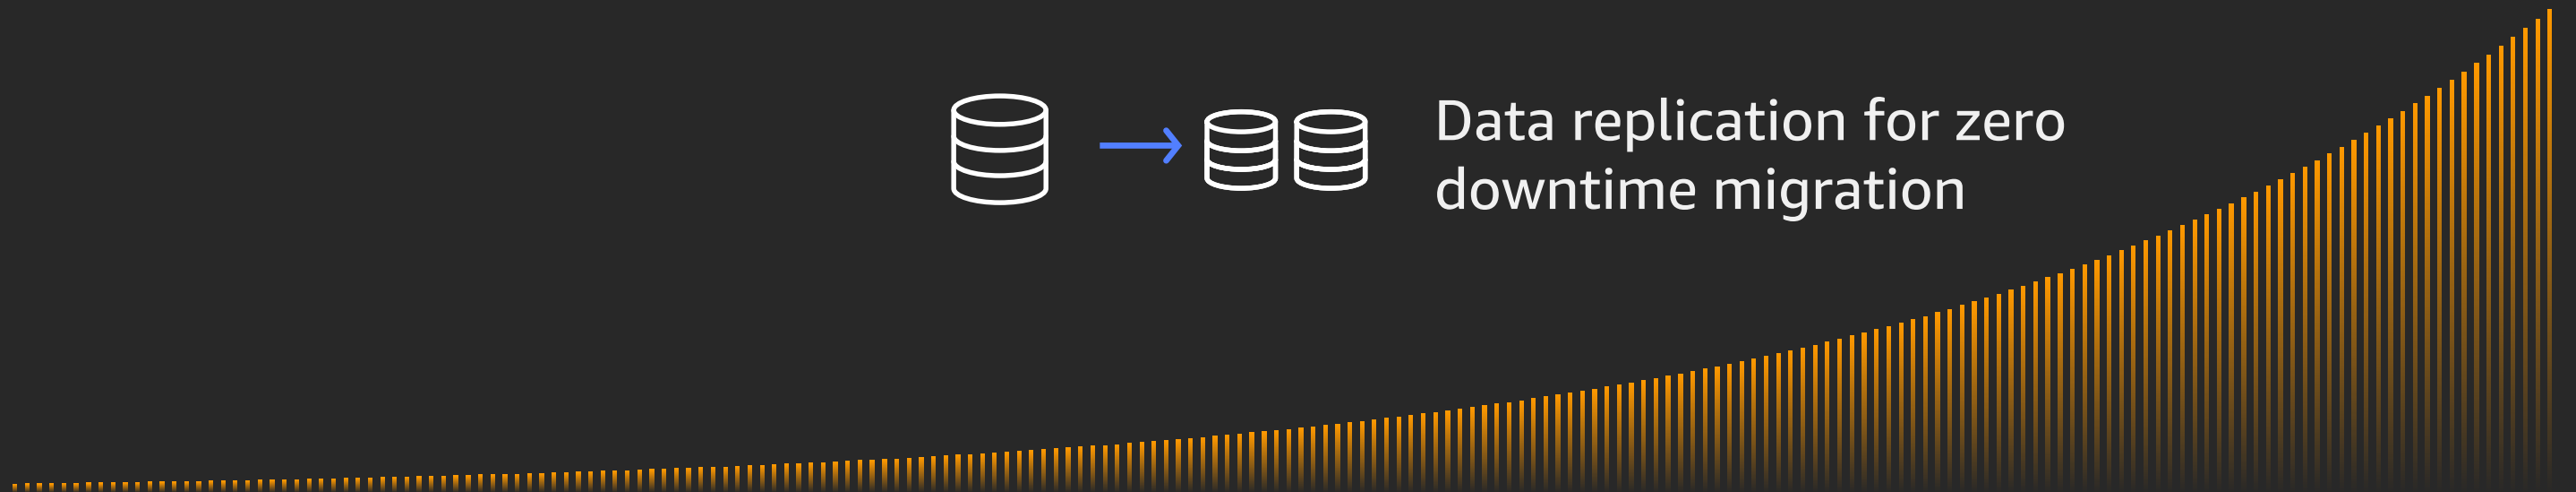
Migrate between databases



Automated schema conversion



Data replication for zero
downtime migration





Legacy data warehouse migration tips

When moving from legacy row-based data warehouses

- Denormalize tables where it makes sense (predicate columns in the fact table)
- Avoid date dimension tables
- Amazon Redshift is efficient with wide tables because of columnar storage and compression

When moving from SMP legacy data warehouses

- Colocation of tables is required for fast JOINS
Leverage DIST STYLE ALL/AUTO or KEY
- Amazon Redshift is designed for big data (>100 GB to PB scale)
Smaller datasets consider Amazon Aurora PostgreSQL
- Wrap workflows in explicit transactions

Leverage Amazon Redshift stored procedures for faster migrations

PL/pgSQL stored procedures were added to make porting legacy procedures easier



Additional resources



AWS Labs on GitHub – Amazon Redshift

<https://github.com/aws-labs/amazon-redshift-utils>

<https://github.com/aws-labs/amazon-redshift-monitoring>

<https://github.com/aws-labs/amazon-redshift-udfs>

Admin scripts

Collection of utilities for running diagnostics on your cluster

Admin views

Collection of utilities for managing your cluster, generating schema DDL, and so on

Analyze Vacuum utility

Utility that you can schedule to vacuum and analyze the tables within your Amazon Redshift cluster

Column Encoding utility

Utility that applies optimal column encoding to an established schema with data already loaded



AWS Big Data Blog – Amazon Redshift

Amazon Redshift Engineering's Advanced Table Design Playbook

<https://aws.amazon.com/blogs/big-data/amazon-redshift-engineerings-advanced-table-design-playbook-preamble-prerequisites-and-prioritization/>

–Zach Christopherson

Top 10 Performance Tuning Techniques for Amazon Redshift

<https://aws.amazon.com/blogs/big-data/top-10-performance-tuning-techniques-for-amazon-redshift/>


–Ian Meyers and Zach Christopherson

Twelve Best Practices for Amazon Redshift Spectrum

<https://aws.amazon.com/blogs/big-data/10-best-practices-for-amazon-redshift-spectrum/>

–Po Hong and Peter Dalton

AWS Training and Certification



Training for the whole team

Explore tailored learning paths for customers and partners



Flexibility to learn your way

Build cloud skills with 550+ free digital training courses, or dive deep with classroom training



Validate skills with AWS Certification

Demonstrate expertise with an industry-recognized credential



Education Programs

Find entry-level cloud talent with AWS Academy and AWS re/Start

aws.amazon.com/training



Thank you!

Aneesh Chandra PN

nanhyama@amazon.com